Bytes file

Created	@Jun 10, 2021 11:17 AM
Tags	

Exploratory Data Analysis

Number of data points in each class

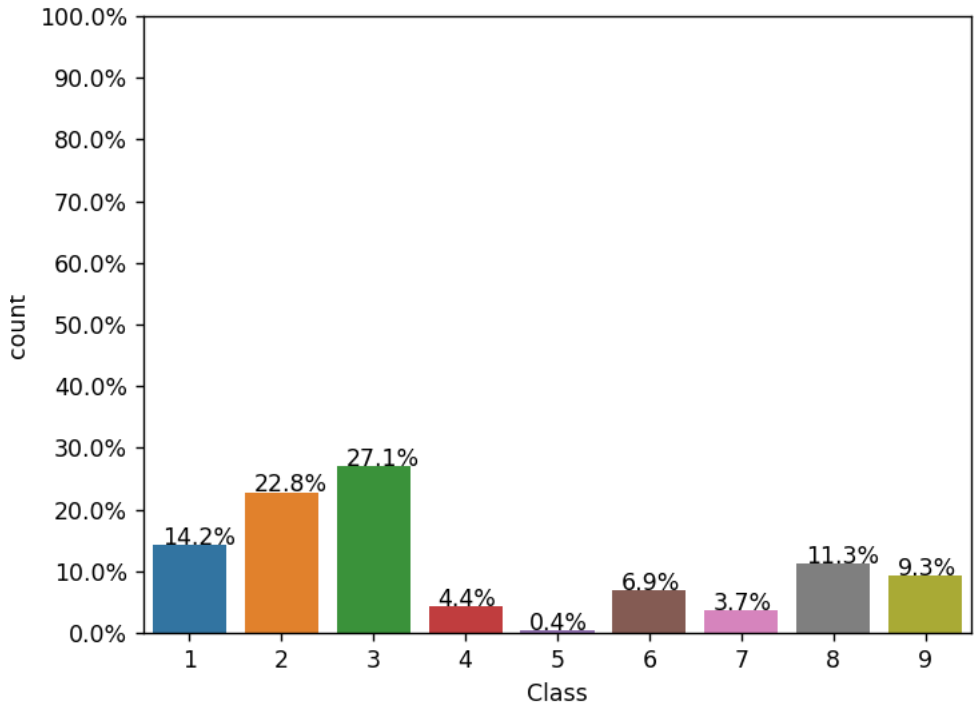

Observation:- Class 5 has less number of data points, **Imbalance data problem.**

File size as feature

ID	File Name	Size
0	01azqd4InC7m9JpocGv5	4.234863
1	01IsoiSMh5gxyDYTI4CB	5.538818
2	01jsnpXSAIgw6aPeDxrU	3.887939
3	01kcPWA9K2BOxQeS5Rju	0.574219
4	01SuzwMJEIXsK7A8dQbl	0.370850

Box plot of file size as feature

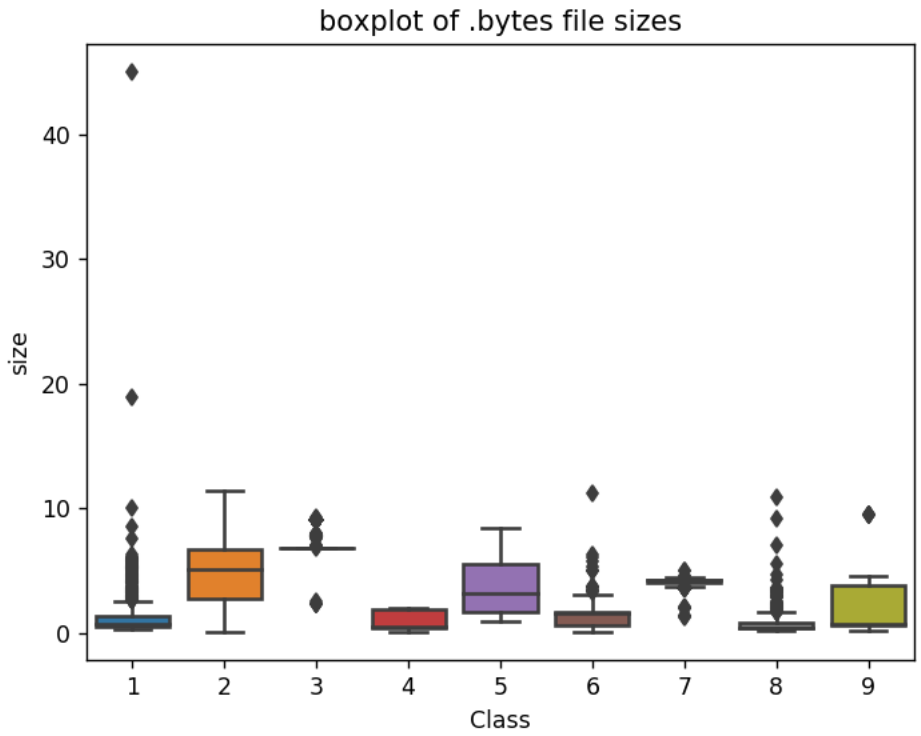

Observation:- Class 2, 5 and 9 can be easily distinguished from other classes, using only the file size feature

Copy of Bag of word as feature of the file

ID	File Name	0	1	2	3	4	5	6	7	8	...	f9	fa	fb	fc
0	01azqd4InC7m9JpocGv5	601905	3905	2816	3832	3345	3242	3650	3201	2965	...	3101	3211	3097	2758
1	01IsoiSMh5gxyDYTI4CB	39755	8337	7249	7186	8663	6844	8420	7589	9291	...	439	281	302	7639
2	01jsnpXSAlgw6aPeDxrU	93506	9542	2568	2438	8925	9330	9007	2342	9107	...	2242	2885	2863	2471
3	01kcPWA9K2BOxQeS5Rju	21091	1213	726	817	1257	625	550	523	1078	...	485	462	516	1133
4	01SuzwMJEIXsK7A8dQbl	19764	710	302	433	559	410	262	249	422	...	350	209	239	653

Copy of Combining Bag of Words and File size as Features

ID	File Name	0	1	2	3	4	5	6	7	8	...	f9	fa	fb	fc
0	01azqd4InC7m9JpocGv5	601905	3905	2816	3832	3345	3242	3650	3201	2965	...	3101	3211	3097	2758
1	01IsoiSMh5gxyDYTI4CB	39755	8337	7249	7186	8663	6844	8420	7589	9291	...	439	281	302	7639
2	01jsnpXSAlgw6aPeDxrU	93506	9542	2568	2438	8925	9330	9007	2342	9107	...	2242	2885	2863	2471
3	01kcPWA9K2BOxQeS5Rju	21091	1213	726	817	1257	625	550	523	1078	...	485	462	516	1133
4	01SuzwMJEIXsK7A8dQbl	19764	710	302	433	559	410	262	249	422	...	350	209	239	653

Copy of Normalizing the Features

ID	File Name	0	1	2	3	4	5	6	7	8
0	01azqd4InC7m9JpocGv5	0.262806	0.005498	0.001567	0.002067	0.002048	0.001835	0.002058	0.002946	0.002638
1	01IsoiSMh5gxyDYTI4CB	0.017358	0.011737	0.004033	0.003876	0.005303	0.003873	0.004747	0.006984	0.008267
2	01jsnpXSAlgw6aPeDxrU	0.040827	0.013434	0.001429	0.001315	0.005464	0.00528	0.005078	0.002155	0.008104
3	01kcPWA9K2BOxQeS5Rju	0.009209	0.001708	0.000404	0.000441	0.00077	0.000354	0.00031	0.000481	0.000959
4	01SuzwMJElXsK7A8dQbl	0.008629	0.001	0.000168	0.000234	0.000342	0.000232	0.000148	0.000229	0.000376

Multivariate analysis of the Features

Perplexity = 50

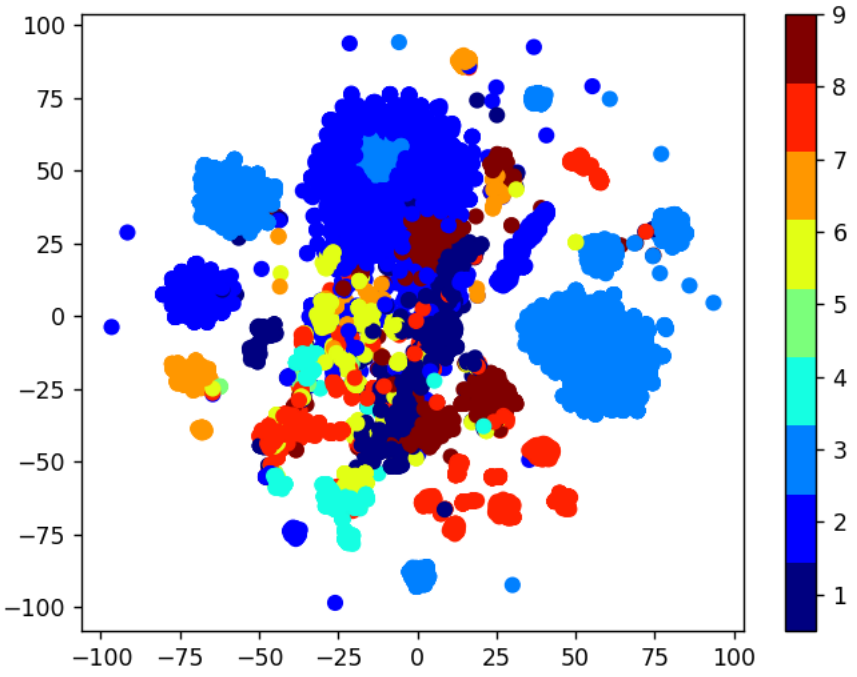

Observation:- Class 2 and 3 are clearly separated whereas other classes have poor distinctions

Perplexity = 30

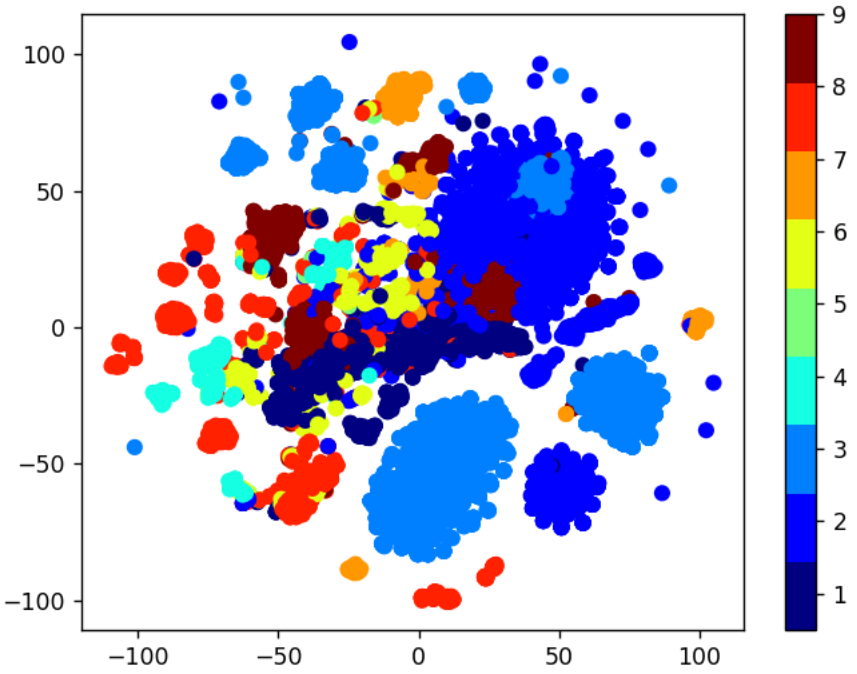

Observation:- Class 2 and 3 are clearly separated whereas other classes have poor distinctions

Test Train Split

Number of data points in train data: 6955
Number of data points in test data: 2174
Number of data points in cross validation data: 1739

Check for distribution of data

We check for the distribution of classes in each split by plotting a histogram.

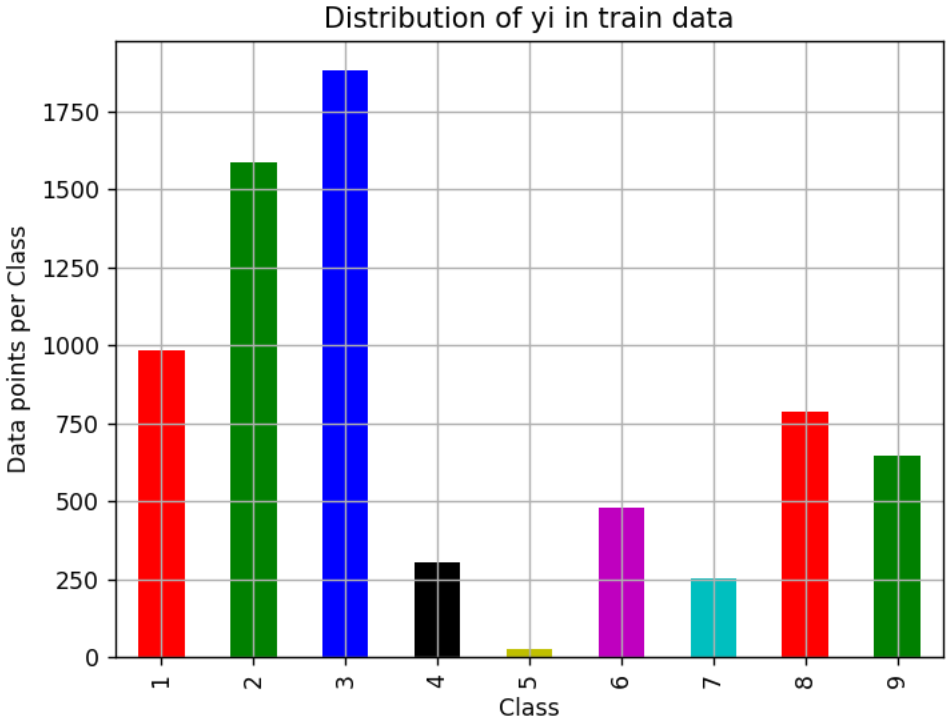
Distribution of yi in train data
Data points per Class
0
250
500
750
1000
1250
1500
1750
1
2
3
4
5
6
7
8
9
Class

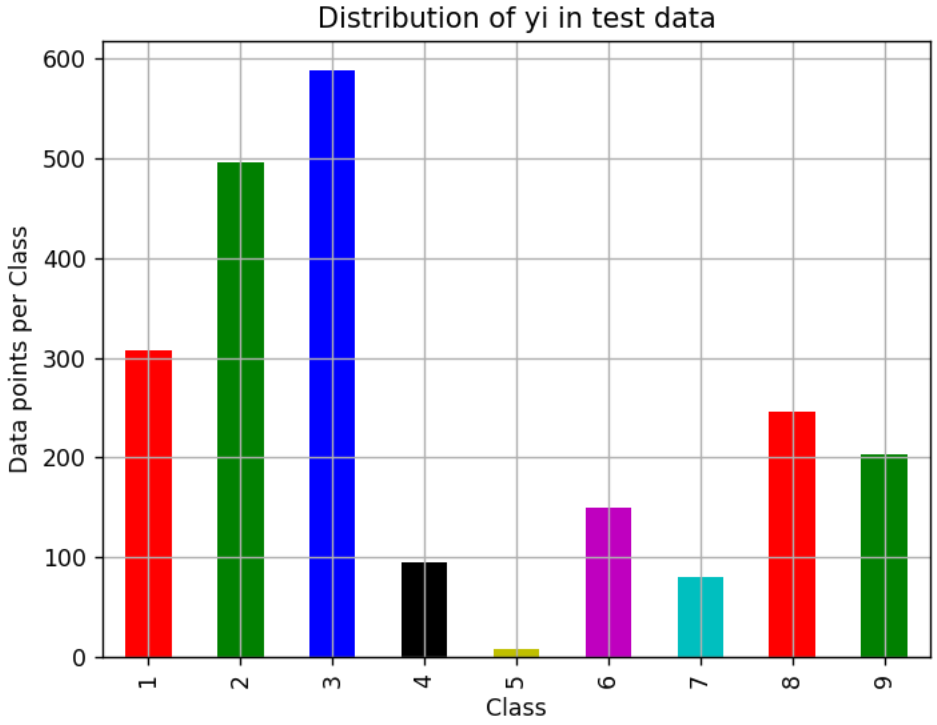
Distribution of yi in test data
Data points per Class
0
100
200
300
400
500
600
1
2
3
4
5
6
7
8
9
Class

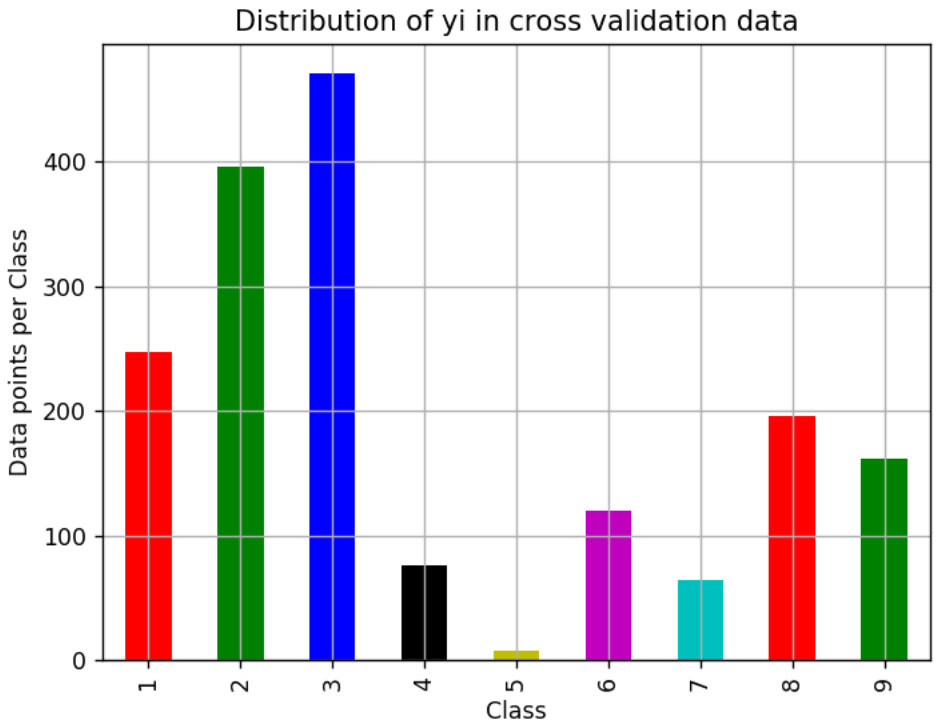

Machine Learning Model

Random Model

Log loss on Cross Validation Data using Random Model 2.46
Log loss on Test Data using Random Model 2.48
Accuracy 11.49

Confusion Matrix

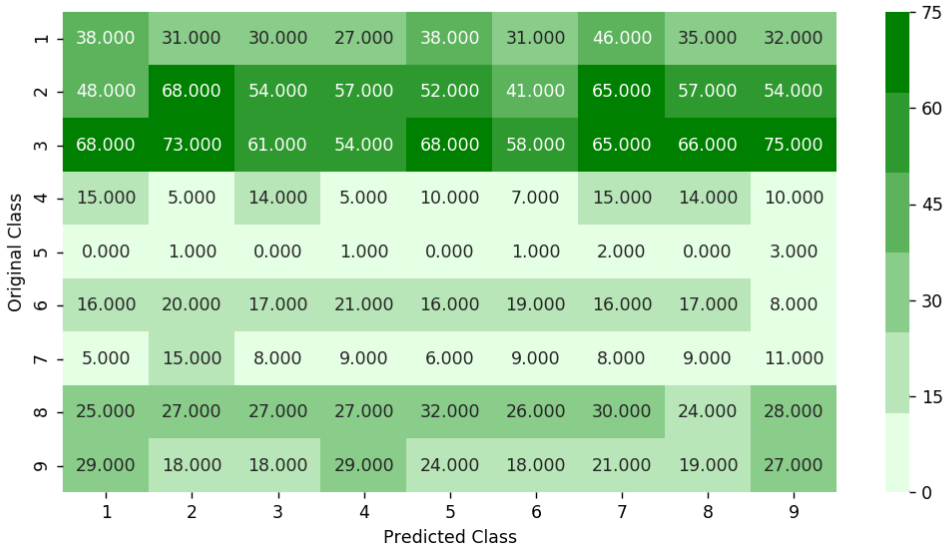

Precision Matrix

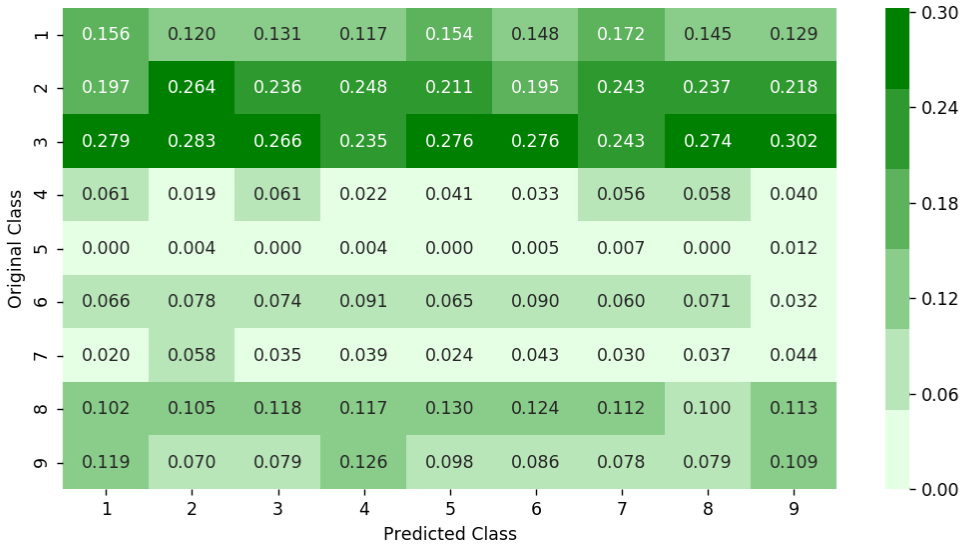

Recall Matrix

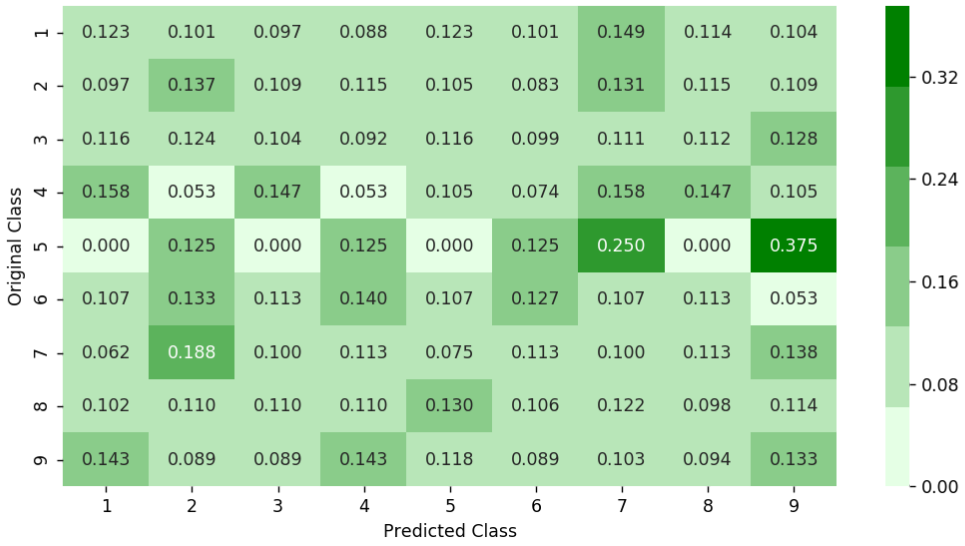

K Nearest Neighbor Classification

Hyperparameter Search

log_loss for k = 1 is 0.225386237304
log_loss for k = 3 is 0.230795229168
log_loss for k = 5 is 0.252421408646
log_loss for k = 7 is 0.273827486888
log_loss for k = 9 is 0.286469181555
log_loss for k = 11 is 0.29623391147
log_loss for k = 13 is 0.307551203154

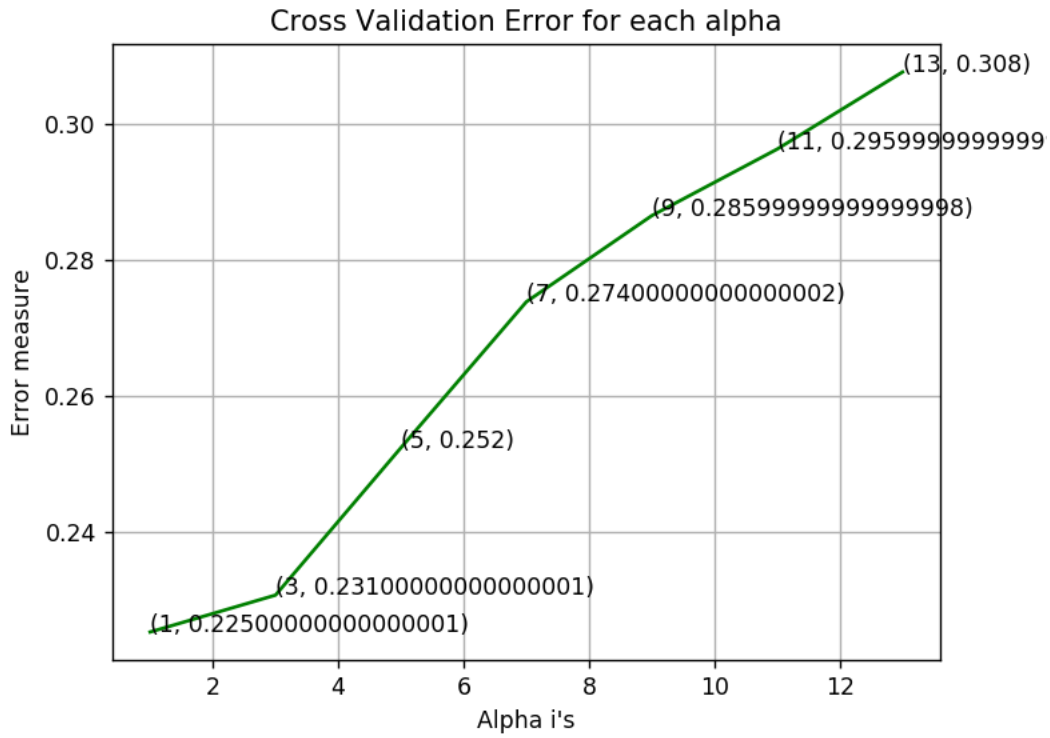

Results from the Best Model

For values of best alpha = 1 The train log loss is: 0.08
For values of best alpha = 1 The cross validation log loss is: 0.23
For values of best alpha = 1 The test log loss is: 0.24
Accuracy 95.49

Confusion Matrix

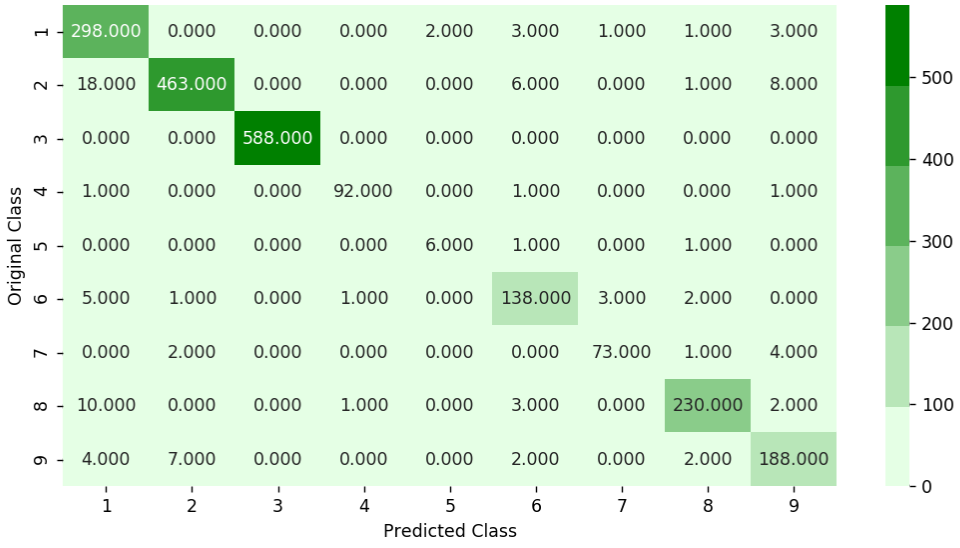

Precision Matrix

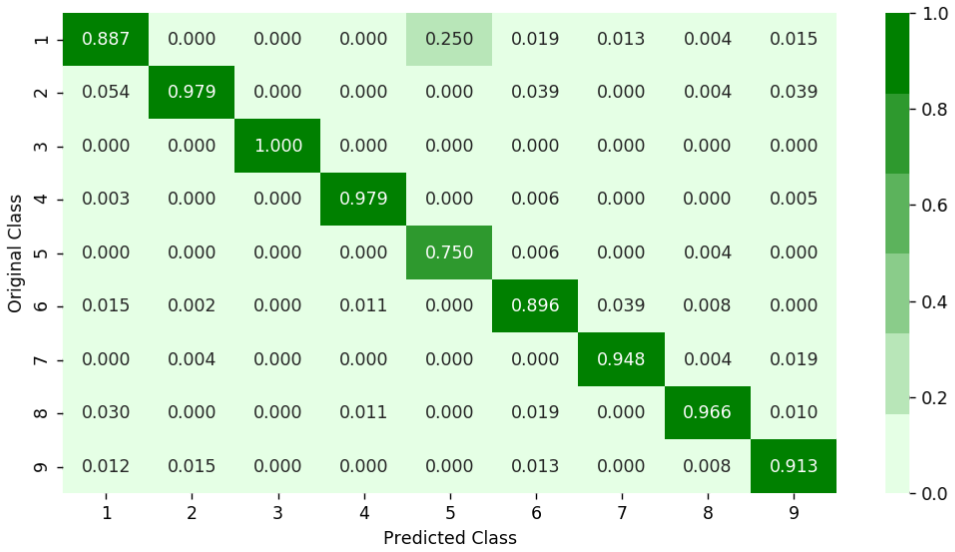

Recall Matrix

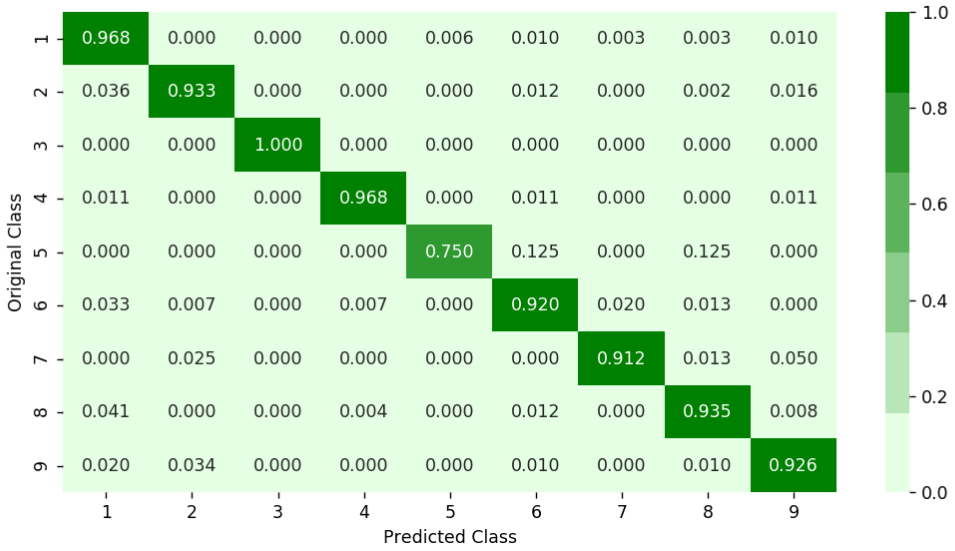

Logistic Regression

Hyperparameter Search

log_loss for c = 1e-05 is 1.56916911178
log_loss for c = 0.0001 is 1.57336384417
log_loss for c = 0.001 is 1.53598598273
log_loss for c = 0.01 is 1.01720972418
log_loss for c = 0.1 is 0.857766083873

log_loss for c = 1 is 0.711154393309
log_loss for c = 10 is 0.583929522635
log_loss for c = 100 is 0.549929846589
log_loss for c = 1000 is 0.624746769121

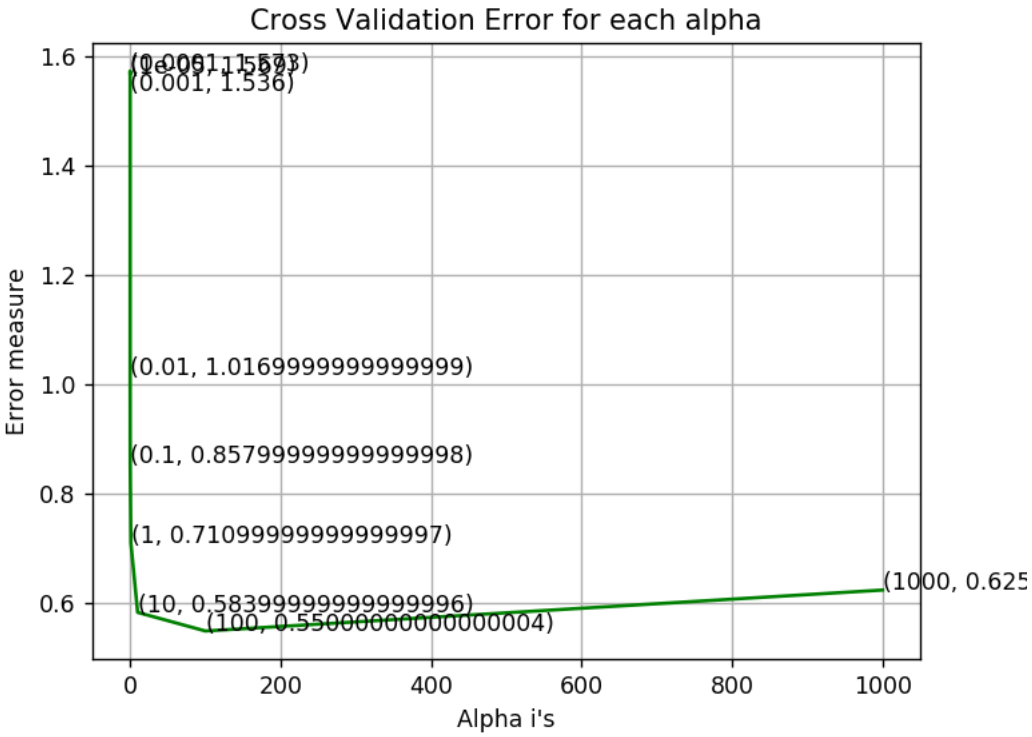

Results from the Best Model

log loss for train data 0.50
log loss for cv data 0.55
log loss for test data 0.53
Number of misclassified points 87.67

Confusion Matrix

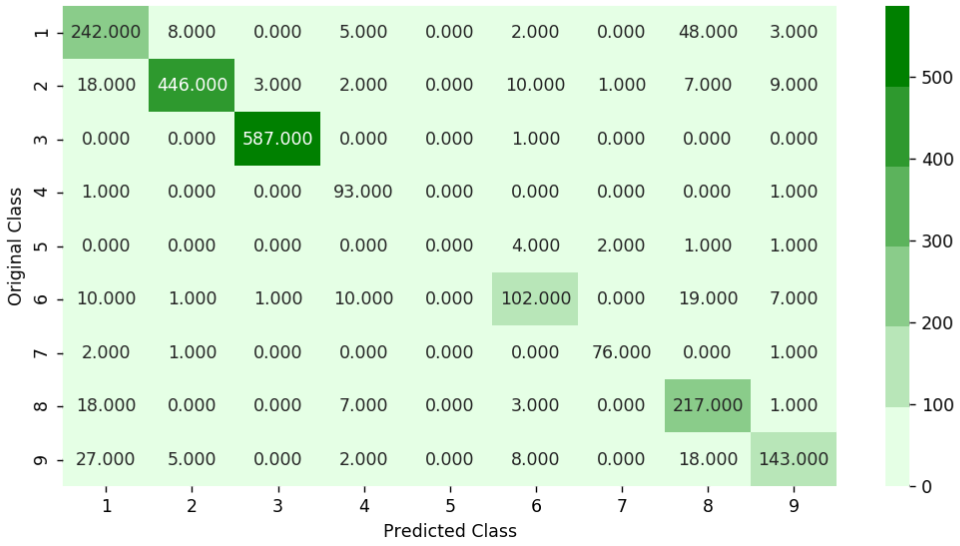

Precision Matrix

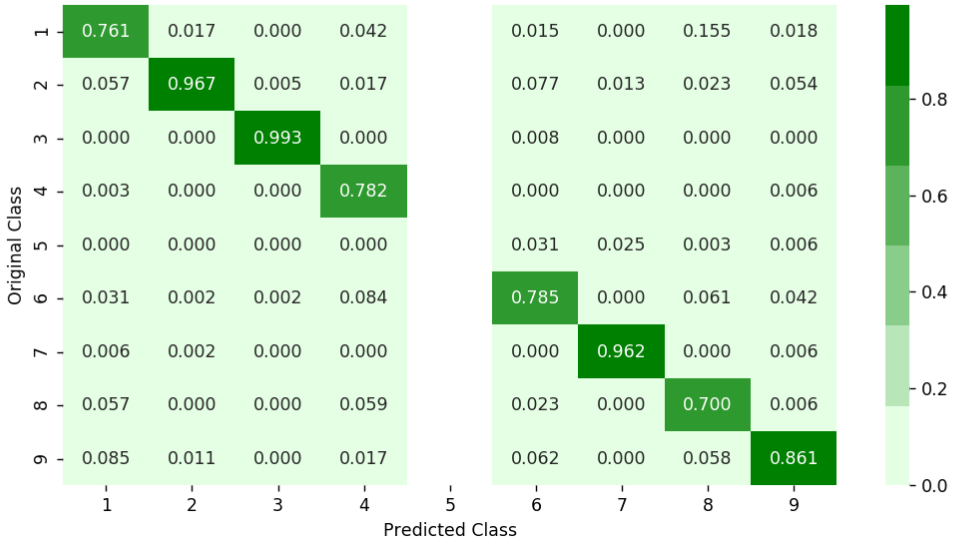

Recall Matrix

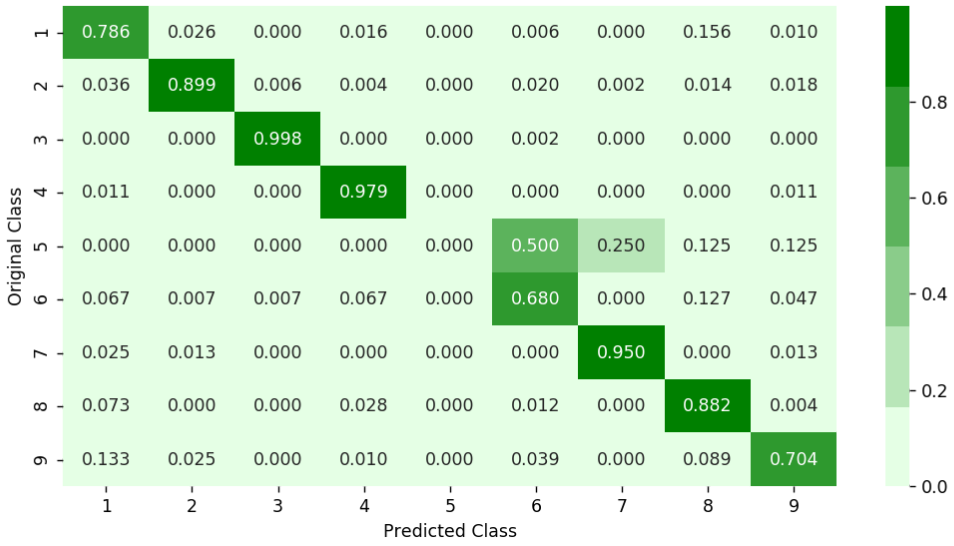

Random Forest Classifier

Hyperparameter Search

log_loss for c = 10 is 0.106357709164
log_loss for c = 50 is 0.0902124124145
log_loss for c = 100 is 0.0895043339776
log_loss for c = 500 is 0.0881420869288
log_loss for c = 1000 is 0.0879849524621
log_loss for c = 2000 is 0.0881566647295
log_loss for c = 3000 is 0.0881318948443

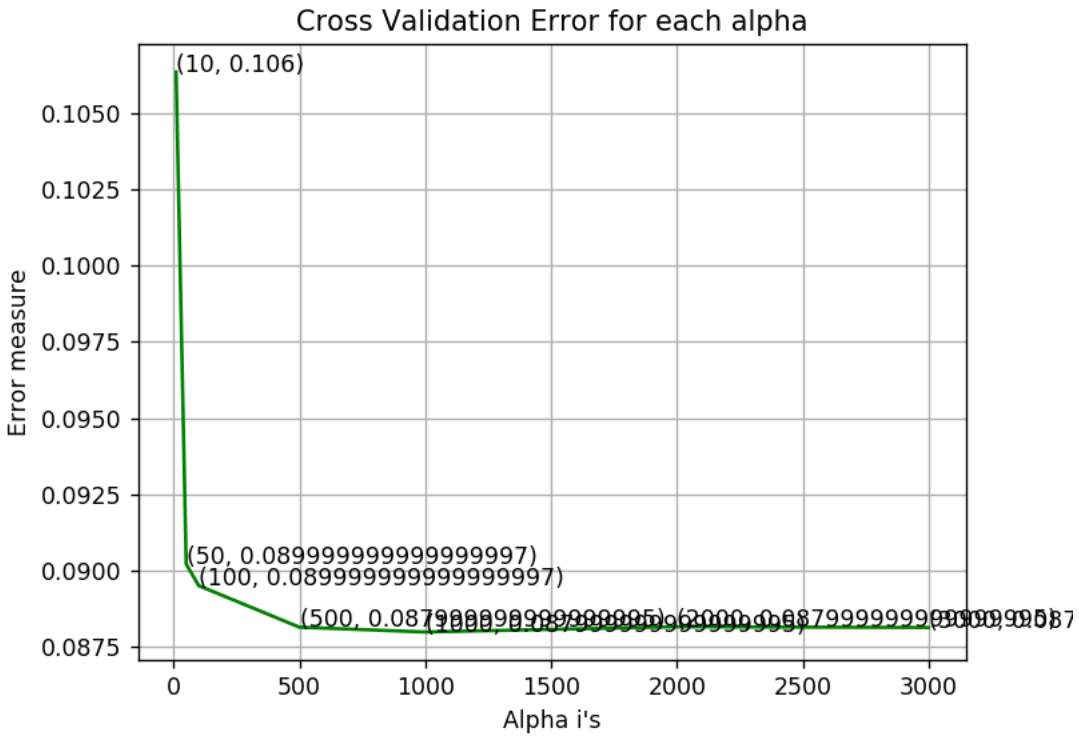

Results from the Best model

For values of best alpha = 1000 The train log loss is: 0.031
For values of best alpha = 1000 The cross validation log loss is: 0.09
For values of best alpha = 1000 The test log loss is: 0.08
Accuracy 96.76

Confusion Matrix

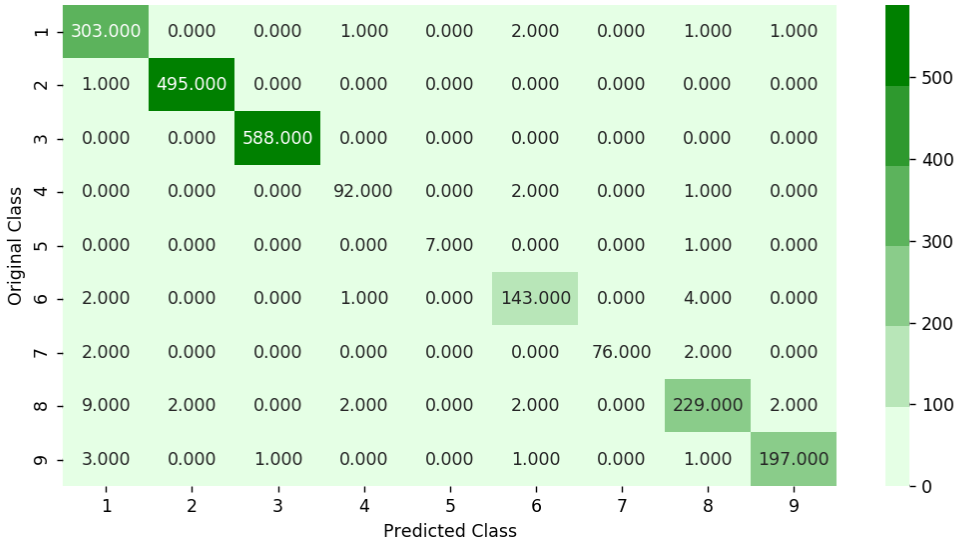

Precision Matrix

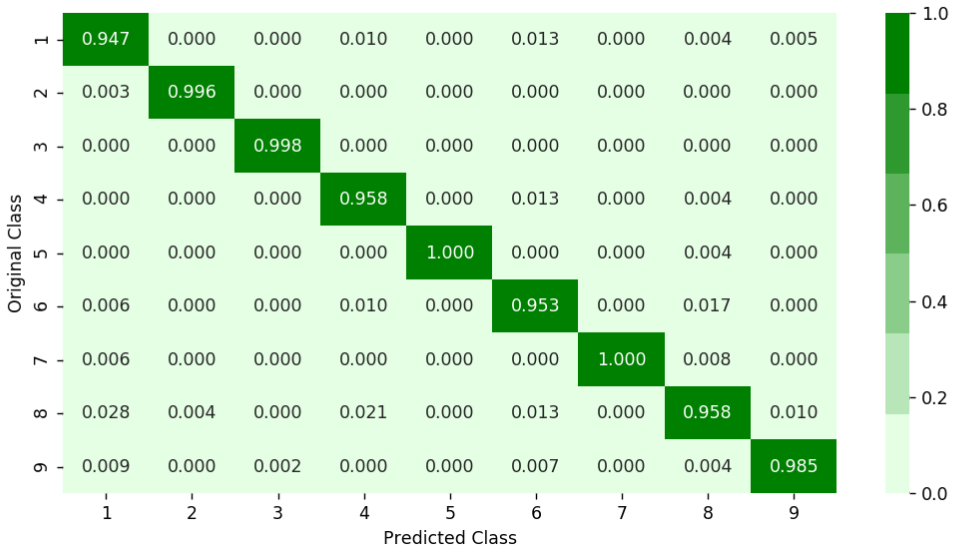

Recall Matrix

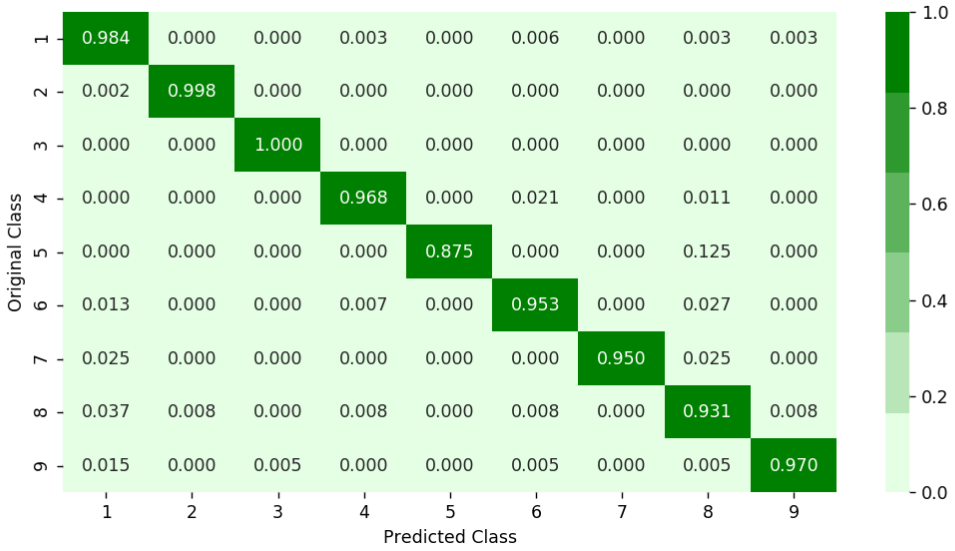

XgBoost Classification

Hyperparameter Search

log_loss for c = 10 is 0.20615980494
log_loss for c = 50 is 0.123888382365
log_loss for c = 100 is 0.099919437112
log_loss for c = 500 is 0.0931035681289

log_loss for c = 1000 is 0.0933084876012
log_loss for c = 2000 is 0.0938395690309

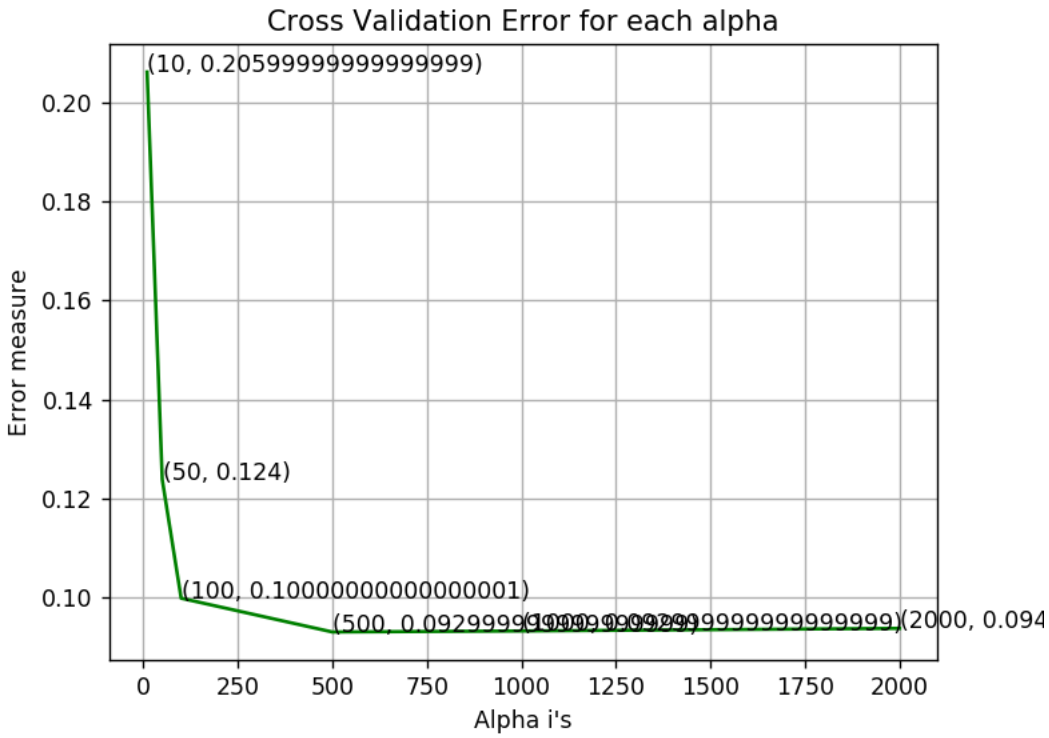

Results from the Best Model

For values of best alpha = 500 The train log loss is: 0.022
For values of best alpha = 500 The cross validation log loss is: 0.09
For values of best alpha = 500 The test log loss is: 0.08
Accuracy 98.67

Confusion Matrix

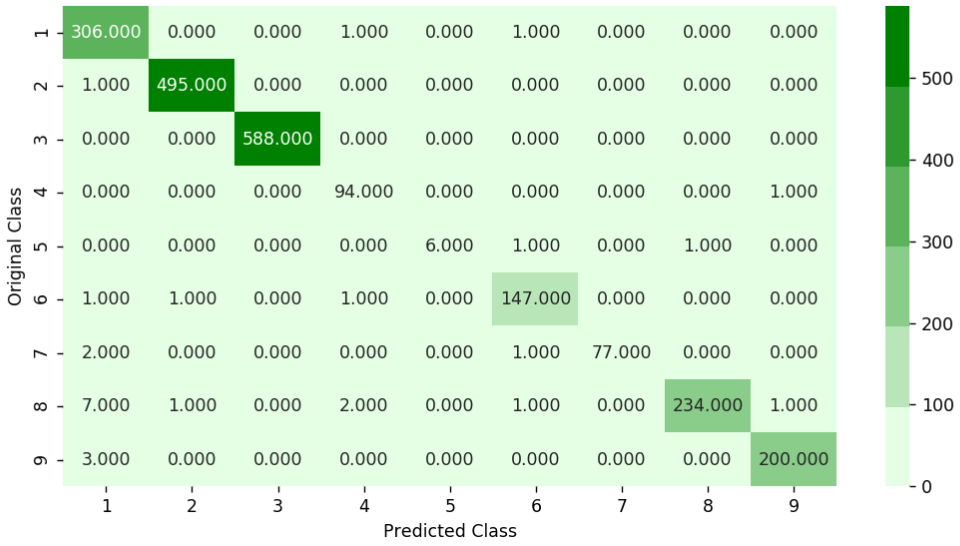

Precision Matrix

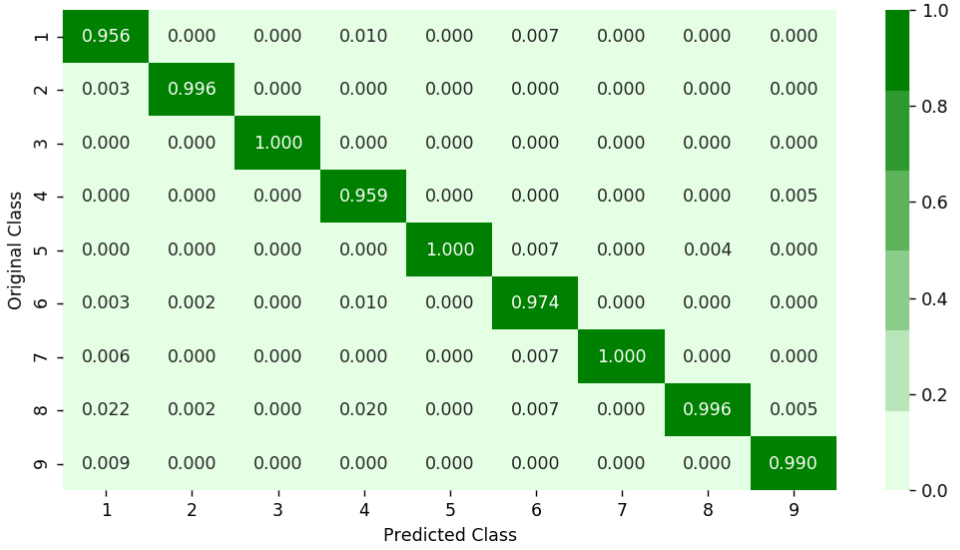

Recall Matrix

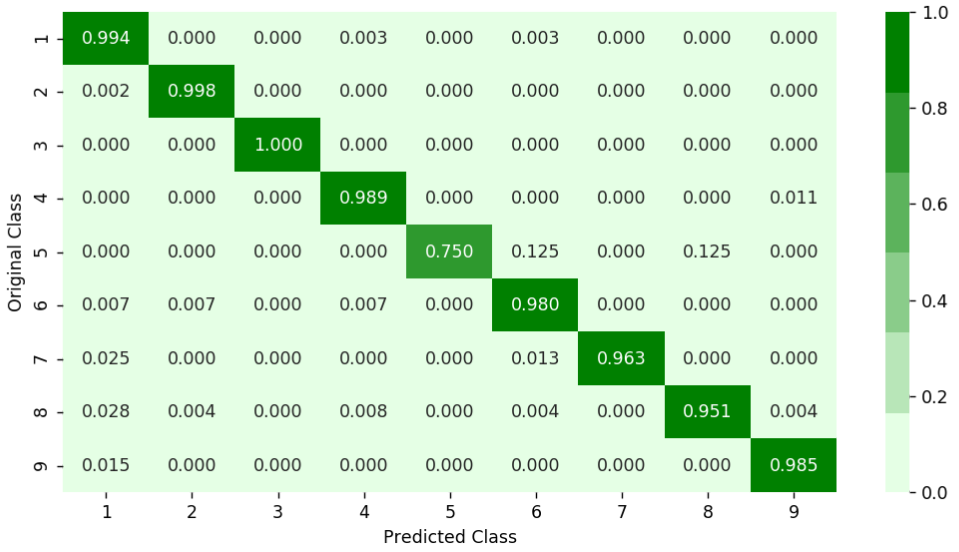

XgBoost Classification with best hyper parameters using Random Search

```
Fitting 3 folds for each of 10 candidates, totalling 30 fits


[Parallel(n_jobs=-1)]: Done   2 tasks      | elapsed:   26.5s
[Parallel(n_jobs=-1)]: Done   9 tasks      | elapsed:  5.8min
[Parallel(n_jobs=-1)]: Done  19 out of  30 | elapsed:  9.3min remaining:  5.4min
[Parallel(n_jobs=-1)]: Done  23 out of  30 | elapsed: 10.1min remaining:  3.1min
[Parallel(n_jobs=-1)]: Done  27 out of  30 | elapsed: 14.0min remaining:  1.6min
[Parallel(n_jobs=-1)]: Done  30 out of  30 | elapsed: 14.2min finished


RandomizedSearchCV(cv=None, error_score='raise',
          estimator=XGBClassifier(base_score=0.5, colsample_bylevel=1, colsample_bytree=1,
       gamma=0, learning_rate=0.1, max_delta_step=0, max_depth=3,
       min_child_weight=1, missing=None, n_estimators=100, nthread=-1,
       objective='binary:logistic', reg_alpha=0, reg_lambda=1,
       scale_pos_weight=1, seed=0, silent=True, subsample=1),
          fit_params=None, iid=True, n_iter=10, n_jobs=-1,
          param_distributions={'learning_rate': [0.01, 0.03, 0.05, 0.1, 0.15, 0.2], 'n_estimators': [100, 200, 500, 1000, 2000], 'max_dep
th': [3, 5, 10], 'colsample_bytree': [0.1, 0.3, 0.5, 1], 'subsample': [0.1, 0.3, 0.5, 1]},
          pre_dispatch='2*n_jobs', random_state=None, refit=True,
          return_train_score=True, scoring=None, verbose=10)
```

Best Parameters

{'subsample': 1, 'n_estimators': 500, 'max_depth': 5, 'learning_rate': 0.05, 'colsample_bytree': 0.5}

Results from the Best Parameter Model

train loss 0.022
cv loss 0.09
test loss 0.08 Accuracy 98.67